

UNIT 4

FL with Differential Privacy

FL with Differential Privacy

- Federated Learning (FL) with **Differential Privacy (DP)** is a technique used to enhance the privacy of decentralized machine learning models.
- It ensures that **individual participants' data remains private** while still contributing to the global model.
- This combination is crucial in domains such as healthcare, finance, and smart grid networks, where data privacy is a priority.

Differential Privacy (DP)

- A mathematical framework that adds controlled noise to data or model updates to prevent the identification of individual data points.
- Ensures that an observer cannot determine whether a specific individual's data was used in training.

How FL and DP Work Together

- **Local DP (LDP):** Noise is added at the client level before sending updates to the server.
 - Each user add noise to their data **before** sending it to the server. This provides **stronger privacy** but often results in **lower utility** due to high noise.
- **Central DP (CDP):** The server applies noise to aggregated updates.
 - Data is collected by a trusted server, and noise is added **centrally** (after collection). It usually provides **better utility** but requires trust in the data collector.
- **Hybrid Approaches:** A combination of LDP and CDP for better privacy-utility trade-offs.

Techniques Used

- **Differentially Private Stochastic Gradient Descent (DP-SGD):** Adds **noise** to gradients during training.
- **Randomized Response:** Clients **randomize their updates** before sending them to the server.
- **Clipping and Noise Addition:** Limits the influence of any single data point and adds Gaussian/Laplace noise.
 - **Why it's used:** Some data points might have **extreme values** that could disproportionately influence the result (e.g., during training).
 - **How it works:** A norm (e.g., **L2 norm**) is applied, and if a data point exceeds a threshold, it's scaled down to the limit.
 - **Example:** In training, if a user's data generates a gradient with **norm 15**, but the **clip norm is set to 5**, the gradient is scaled down proportionally so its norm becomes 5.

Randomize Response

- Instead of directly sharing a sensitive value, a user applies a probabilistic mechanism to randomize their response before sending it.
- This makes it difficult for an observer (such as a server in FL) to determine whether a specific response is the true value or a randomized one.

Randomize Response

1. **Basic Example (Binary Response)** Suppose a survey asks, "*Have you ever committed tax fraud?*" Instead of answering truthfully, the participant:

- Flips a coin.
- If **heads**, they answer truthfully.
- If **tails**, they flip a second coin:
 - If **heads**, they say "Yes."
 - If **tails**, they say "No."

This mechanism ensures plausible deniability while still allowing researchers to estimate the proportion of actual "Yes" responses in a population.

Randomize Response

Mathematical Formulation If a true value $x \in \{0, 1\}$ (e.g., "No" = 0, "Yes" = 1) is reported with probability p and randomized with probability $1 - p$, then:

$$P(\hat{x} = 1) = p \cdot x + (1 - p) \cdot 0.5$$

where \hat{x} is the reported response.

This allows the system to statistically estimate the true distribution while protecting individual privacy.

Where:

- \hat{x} is the reported (randomized) value.
- $x \in \{0, 1\}$ is the true value (e.g., "No" = 0, "Yes" = 1).
- p is the probability of reporting the true value.
- $(1 - p)$ is the probability of randomizing the response.
- The term 0.5 accounts for equal probability of flipping the response when randomization happens.

Randomize Response

Example Calculation

If:

- $p = 0.8$ (80% chance of reporting truthfully),
- $x = 1$ (true answer is Yes),
- $q = 0.5$ (random response is equally likely to be 0 or 1),

Then:

$$P(\hat{x} = 1) = 0.8 \cdot 1 + (1 - 0.8) \cdot 0.5$$

$$P(\hat{x} = 1) = 0.8 + 0.1 = 0.9$$

Thus, there is a **90% probability** that the reported answer will be "Yes" (even though true responses were randomized in some cases).

Randomize Response (Challenges)

- **Accuracy Loss:** Excessive noise can degrade model performance.
- **Parameter Tuning Required:** The probability p must be carefully chosen to balance privacy and utility.
- **Aggregation Complexity:** Requires statistical post-processing to reconstruct meaningful insights from noisy data.

Advantages of Randomized Response

- ✓ **Strong Privacy Protection:** Ensures individual data points remain hidden.
- ✓ **Simple and Efficient:** Does not require complex cryptographic computations like homomorphic encryption.
- ✓ **Useful in Decentralized Settings:** Can be applied in federated learning where no trusted third party exists.

Clipping (Gradient Clipping)

Mathematical Formulation

Given a local model update g (a gradient vector) computed by a client, we **clip** it using a pre-defined threshold C :

$$g' = \frac{g}{\max(1, \frac{\|g\|}{C})}$$

Where:

- g is the original gradient,
- C is the clipping norm,
- g' is the clipped gradient.

Why is Clipping Important?

- ✓ **Prevents Outliers:** Avoids any single client contributing an extremely large gradient, which could reveal sensitive data.
- ✓ **Improves Training Stability:** Prevents gradient explosion, leading to more stable convergence.
- ✓ **Enables Differential Privacy:** Clipping ensures that all gradients have a bounded impact before noise is added.

Clipping (Gradient Clipping)

What is Noise Addition?

After clipping, **random noise** is added to the gradients before they are sent to the central server.

This ensures **differential privacy (DP)** by making it hard for an adversary to infer information about individual training samples.

Mathematical Formulation

The **noisy gradient** \tilde{g} is obtained by adding noise to the clipped gradient:

$$\tilde{g} = g' + \mathcal{N}(0, \sigma^2 C^2)$$

Where:

- g' is the clipped gradient,
- $\mathcal{N}(0, \sigma^2 C^2)$ is Gaussian noise with **variance** $\sigma^2 C^2$,
- σ controls the noise level (higher $\sigma \rightarrow$ stronger privacy, lower accuracy).

Putting It Together: Clipping + Noise Addition in Federated Learning

1. Clients compute local gradients based on their private data.
2. Gradient clipping is applied to bound the maximum contribution of any client.
3. Gaussian noise is added to ensure differential privacy.
4. Noisy, clipped gradients are sent to the central server.
5. The server aggregates gradients and updates the global model.

Challenges & Trade-offs

● Privacy vs. Utility Trade-off

- More noise (higher σ) → Better privacy, but lower model accuracy.
- Less noise (lower σ) → Higher accuracy, but weaker privacy.

● Choosing an Appropriate Clipping Norm (C)

- Too high: Privacy is weaker since gradients are not well-bounded.
- Too low: Useful information is lost, degrading model performance.

● Computational Overhead

- Noise addition and clipping introduce additional computations, especially in large-scale federated learning deployments.

Secure Aggregation Techniques

- **Homomorphic Encryption (HE):** Allows computations on encrypted data without decryption, enabling privacy-preserving aggregation in federated learning and cloud environments.
- **Differential Privacy (DP):** Introduces controlled noise to aggregated data to prevent individual data points from being inferred.
- **Secure Multi-Party Computation (SMPC):** Enables multiple parties to jointly compute a function while keeping their inputs private.
- **Secret Sharing:** Data is split into multiple shares distributed among participants, ensuring no single entity can access complete information.
- **Verifiable Computation:** Ensures that aggregated results are correct and tamper-proof using cryptographic proofs.

Encryption Techniques

- **AES (Advanced Encryption Standard):** A widely used symmetric encryption algorithm ensuring data confidentiality.
- **RSA (Rivest-Shamir-Adleman):** An asymmetric encryption technique used for secure key exchange and encryption.
- **Elliptic Curve Cryptography (ECC):** Provides strong security with smaller key sizes, making it efficient for IoT and mobile devices.
- **Fully Homomorphic Encryption (FHE):** Allows complex computations on encrypted data without decryption, ideal for privacy-sensitive applications.
- **Attribute-Based Encryption (ABE):** Encrypts data based on user attributes, providing fine-grained access control.

Homomorphic Encryption

- Homomorphic Encryption (HE) is a form of encryption that allows computations to be **performed on ciphertexts** without decrypting them.
- The results, when decrypted, match those obtained if operations were performed on the plaintext.
- This property is particularly useful in privacy-preserving computations, such as secure cloud computing, encrypted machine learning, and confidential data analysis.

Types of Homomorphic Encryption

Partially Homomorphic Encryption (PHE):

- Supports only a single type of operation (either addition or multiplication).
- Example: RSA and ElGamal encryption support multiplication, while Paillier encryption supports addition.

Somewhat Homomorphic Encryption (SWHE):

- Supports a limited number of both addition and multiplication operations.
- Constraint: The number of operations is restricted due to noise accumulation.

Fully Homomorphic Encryption (FHE):

- Supports an unlimited number of additions and multiplications.
- Enables complex computations on encrypted data without decryption.
- First practical FHE scheme was introduced by Craig Gentry in 2009.

Types of Homomorphic Encryption

Secure Cloud Computing: Users can process encrypted data in the cloud without exposing sensitive information.

Privacy-Preserving Machine Learning: Enables training models on encrypted data without revealing the raw dataset.

Secure Electronic Voting: Ensures confidentiality while allowing mathematical operations to count votes.

Financial and Healthcare Data Security: Protects sensitive records while enabling computations on them.

Blockchain & Secure Transactions: Ensures privacy in smart contracts and decentralized finance (DeFi) applications.

Challenges of Homomorphic Encryption

Computational Overhead: HE operations are significantly slower than plaintext computations.

Key Management: Requires complex key management for security.

Noise Accumulation: Noise grows with each operation, requiring bootstrapping (refreshing ciphertexts).

Key Properties of PHE

- **Single Operation Support:** PHE allows either addition or multiplication but not both.
- **Computational Efficiency:** Compared to FHE, PHE is more efficient since it doesn't require complex bootstrapping techniques.
- **Security:** PHE schemes are based on hard mathematical problems like integer factorization or discrete logarithms, making them secure under well-known cryptographic assumptions.

PHE

1. Paillier Cryptosystem (Addition Homomorphism)

- Given two encrypted values $Enc(m_1)$ and $Enc(m_2)$, the system allows:

$$Enc(m_1) \cdot Enc(m_2) = Enc(m_1 + m_2)$$

- This means that multiplying encrypted values results in the encryption of their sum.

2. RSA Cryptosystem (Multiplication Homomorphism)

- Given two encrypted values $Enc(m_1)$ and $Enc(m_2)$, the system allows:

$$Enc(m_1) \times Enc(m_2) = Enc(m_1 \cdot m_2)$$

- This means that multiplying ciphertexts results in the encryption of their product.

Applications of PHE

- **Secure Online Voting** – Ensures that votes remain private while allowing addition-based tallying.
- **Privacy-Preserving Finance** – Banks can perform encrypted transactions without accessing user data.
- **Cloud Computing** – Enables cloud providers to process encrypted data without decrypting it.
- **Digital Rights Management (DRM)** – Helps in controlled access and manipulation of encrypted content.

Bootstrapping

- In FHE schemes, every homomorphic operation (addition or multiplication) introduces noise into the ciphertext.
- If too many operations are performed, the noise grows beyond a threshold, making decryption fail.
- **Bootstrapping resets this noise level**, effectively "refreshing" the ciphertext to allow further computations.

How Bootstrapping Works

1. Encrypt the Encrypted Ciphertext

- The noisy ciphertext is re-encrypted using a special procedure.

2. Evaluate a Decryption Circuit Homomorphically

- The system performs a decryption operation on the encrypted ciphertext **while it is still encrypted**.

3. Output a Fresh, Less Noisy Ciphertext

- The resulting ciphertext contains the same message but with reduced noise, allowing continued operations.

Bootstrapping Challenges

- **Computationally Expensive** – Bootstrapping requires executing a complex decryption function homomorphically, which is slow.
- **High Latency** – It takes significant processing power and time, making FHE impractical for real-time applications.
- **Large Key and Ciphertext Size** – Additional complexity leads to larger memory and storage requirements.

Paillier Cryptosystem (Additive Homomorphism)

The Paillier cryptosystem supports homomorphic **addition** of encrypted values.

Key Generation:

1. Choose two large prime numbers p and q .
2. Compute $n = p \cdot q$ and $\lambda = \text{lcm}(p - 1, q - 1)$.
3. Select a generator g such that $g^\lambda \pmod{n^2}$ is invertible.
4. Compute the decryption function $\mu = (\text{L}(g^\lambda \pmod{n^2}))^{-1} \pmod{n}$, where $\text{L}(x) = \frac{x-1}{n}$.
5. **Public Key:** (n, g) , **Private Key:** (λ, μ) .

Encryption:

For a message $m \in \mathbb{Z}_n$, select a random number $r \in \mathbb{Z}_n^*$ and compute:

$$Enc(m) = g^m \cdot r^n \mod n^2$$

Decryption:

Given ciphertext c :

$$Dec(c) = L(c^\lambda \mod n^2) \cdot \mu \mod n$$

Homomorphic Property (Addition):

For two messages m_1, m_2 :

$$\begin{aligned} Enc(m_1) \cdot Enc(m_2) &= (g^{m_1} r_1^n) \cdot (g^{m_2} r_2^n) \mod n^2 \\ &= g^{m_1+m_2} (r_1 r_2)^n \mod n^2 \end{aligned}$$

Since $(r_1 r_2)^n$ is still a random value,

$$Enc(m_1) \cdot Enc(m_2) = Enc(m_1 + m_2)$$

Thus, Paillier encryption **preserves addition** under encryption.

Example

1. Choose two prime numbers:

$$p = 7, \quad q = 11$$

2. Compute n and n^2 :

$$n = p \cdot q = 7 \times 11 = 77$$

$$n^2 = 77^2 = 5929$$

3. Compute $\lambda(n)$:

$$\lambda = \text{lcm}(p - 1, q - 1) = \text{lcm}(6, 10) = 30$$

4. Choose g :

The standard choice is:

$$g = n + 1 = 77 + 1 = 78$$

Example

5. Compute μ :

First, compute:

$$L(g^\lambda \mod n^2) = \frac{g^\lambda - 1}{n}$$

$$78^{30} \mod 5929 = 4631$$

$$L(4631) = \frac{4631 - 1}{77} = 60$$

Compute μ as the modular inverse of 60 mod 77:

$$\mu = 60^{-1} \mod 77 = 38$$

Public Key: $(n, g) = (77, 78)$

Private Key: $(\lambda, \mu) = (30, 38)$

Example

Step 2: Encryption of a Message

Let's encrypt $m = 20$.

1. Choose a random number r :

Let $r = 3$.

2. Compute ciphertext c :

$$c = g^m \cdot r^n \pmod{n^2}$$

$$c = 78^{20} \cdot 3^{77} \pmod{5929}$$

Breaking it down:

$$78^{20} \pmod{5929} = 2087$$

$$3^{77} \pmod{5929} = 1713$$

$$c = (2087 \times 1713) \pmod{5929}$$

$$c = 3573231 \pmod{5929} = 1222$$

So, the encrypted ciphertext is:

$$c = 1222$$

Example

Step 3: Homomorphic Addition

Let's encrypt another message $m' = 15$ and compute the sum homomorphically.

1. Encrypt $m' = 15$ with another random $r' = 5$:

$$c' = g^{15} \cdot 5^{77} \pmod{5929}$$

Assume:

$$c' = 3210$$

2. Compute **homomorphic addition**:

$$c_{\text{sum}} = c \cdot c' \pmod{n^2}$$

$$\begin{aligned} c_{\text{sum}} &= (1222 \times 3210) \pmod{5929} \\ &= 3928620 \pmod{5929} = 2513 \end{aligned}$$

The encrypted sum $Enc(20 + 15) = Enc(35)$ is:

$$c_{\text{sum}} = 2513$$

Example

Step 4: Decryption

Now, let's decrypt $c_{\text{sum}} = 2513$.

1. Compute:

$$L(c_{\text{sum}}^{\lambda} \mod n^2)$$

$$L(2513^{30} \mod 5929)$$

$$2513^{30} \mod 5929 = 2693$$

2. Compute $L(2693)$:

$$L(2693) = \frac{2693 - 1}{77} = 35$$

3. Multiply by μ :

$$m = L(2693) \cdot \mu \mod n$$

$$m = 35 \times 38 \mod 77$$

$$= 1330 \mod 77 = 35$$

Thus, the decrypted message is **35**, which matches **20 + 15**, verifying the homomorphic property.

RSA Cryptosystem (Multiplicative Homomorphism)

The **RSA cryptosystem** supports homomorphic **multiplication** of encrypted values.

Key Generation:

1. Choose two large prime numbers p and q .
2. Compute $n = p \cdot q$ and $\phi(n) = (p - 1)(q - 1)$.
3. Select a public exponent e such that $\gcd(e, \phi(n)) = 1$.
4. Compute the private exponent d such that $e \cdot d \equiv 1 \pmod{\phi(n)}$.
5. **Public Key:** (n, e) , **Private Key:** (d) .

RSA Cryptosystem (Multiplicative Homomorphism)

Encryption:

For a message $m \in \mathbb{Z}_n$, compute:

$$Enc(m) = m^e \pmod{n}$$

Decryption:

$$Dec(c) = c^d \pmod{n} = (m^e)^d \pmod{n} = m^{ed} \pmod{n} = m$$

Homomorphic Property (Multiplication):

For two messages m_1, m_2 :

$$\begin{aligned} Enc(m_1) \cdot Enc(m_2) &= (m_1^e \pmod{n}) \cdot (m_2^e \pmod{n}) \pmod{n} \\ &= (m_1 m_2)^e \pmod{n} = Enc(m_1 \cdot m_2) \end{aligned}$$

Thus, RSA encryption preserves multiplication under encryption.

Secret Sharing

- Secret sharing is a crucial technique used in **secure aggregation** for **federated learning (FL)** to ensure **privacy-preserving model training**.
- Secure aggregation allows multiple clients to collaboratively train a model while keeping their individual model updates private.
- Secret sharing enables a client to split its **model updates (gradients/weights)** into multiple **shares** and distribute them among other participating clients.

Secret Sharing

These shares are structured such that:

- No single client (or a subset of clients below a threshold) can reconstruct the original update.
- Only the **server** (aggregator) can retrieve the sum of all updates without learning individual contributions. (The goal is to allow the **server (aggregator)** to compute the **sum of all local model updates** while ensuring that **individual client contributions remain hidden**.)

How Does This Work?

1. Clients Secret-Share Their Updates

- Each client splits its local update x_i into **multiple shares** using **secret sharing** (e.g., **additive secret sharing** or **Shamir's Secret Sharing**).
- The shares are distributed among other clients.

2. Clients Exchange Shares

- Each client receives shares from other clients but cannot reconstruct any individual update.
- The clients securely compute the sum of the shares and send them to the server.

3. Server Aggregates the Updates

- The server only receives the **sum of all shares**, which is equivalent to the sum of the original updates:

$$\sum_{i=1}^N x_i$$

- Since the individual updates are secret-shared, the server cannot learn **any single client's update**.

Why is the Server Unable to Learn Individual Updates?

- **Lack of Access to Individual Shares:**

The server does not receive individual shares—only the final sum.

- **Randomization:**

Each client's update is **masked** using randomness, making it impossible to infer individual values.

- **Threshold-Based Security:**

If a **threshold secret sharing scheme** (e.g., Shamir's Secret Sharing) is used, the server would need at least **t** out of **N** shares to reconstruct any single update—making privacy robust.

Secret Sharing Techniques

Shamir's Secret Sharing (SSS)

- Based on polynomial interpolation.
- A client splits its update into **n shares** using a polynomial of degree **t-1** (where t is the threshold).
- Only **t or more** clients can reconstruct the original update.
- **Advantages:** Strong theoretical security, protects against collusion.
- **Challenges:** High computational overhead, requires exact threshold reconstruction.

Secret Sharing Techniques

SSS is based on **Lagrange interpolation** over a finite field F. The secret is hidden within a **random polynomial**, and only a sufficient number of shares can reconstruct it.

- **Steps in Shamir's Secret Sharing:**

1. **Secret Definition:**

- Let the secret be a numerical value s (e.g., a model update in FL).
- Choose a **threshold t** and the **total number of shares n** , where $t \leq n$.

Secret Sharing Techniques

2. Polynomial Generation:

- The client creates a random polynomial of degree $t - 1$:

$$f(x) = s + a_1x + a_2x^2 + \dots + a_{t-1}x^{t-1}$$

where a_1, a_2, \dots, a_{t-1} are random coefficients.

3. Share Distribution:

- The client computes n shares:

$$\text{Share}_i = f(x_i) \quad \text{for } i = 1, 2, \dots, n$$

- Each participating client C_i receives a unique share $(x_i, f(x_i))$.

Secret Sharing Techniques

4. Aggregation Phase (At the Server):

- The server receives the **sum of all shares** from participating clients.
- It applies **Lagrange interpolation** to reconstruct the sum of the original model updates:

$$s = f(0) = \sum \lambda_i f(x_i)$$

- Since the individual shares are masked by the polynomial, the server cannot learn individual updates.

Example of Secret Sharing Techniques

We want to share a secret $S = 42$ among **5 participants**, requiring at least **3 shares** to reconstruct it.

We use a **random polynomial** of degree $t - 1$ (where t is the threshold). Since we require **3 shares to reconstruct the secret**, we choose a **quadratic polynomial** (degree 2):

$$f(x) = a_0 + a_1x + a_2x^2$$

where:

- $a_0 = 42$ (the secret)
- a_1 and a_2 are random coefficients (let's pick $a_1 = 7$ and $a_2 = 3$)

Thus, our polynomial is:

$$f(x) = 42 + 7x + 3x^2$$

Example of Secret Sharing Techniques

Step 1: Generate Shares

We compute $f(x)$ for different values of x (choosing $x = 1, 2, 3, 4, 5$):

x	$f(x) = 42 + 7x + 3x^2$	Share $(x, f(x))$
1	$42 + 7(1) + 3(1)^2 = 42 + 7 + 3 = 52$	(1, 52)
2	$42 + 7(2) + 3(2)^2 = 42 + 14 + 12 = 68$	(2, 68)
3	$42 + 7(3) + 3(3)^2 = 42 + 21 + 27 = 90$	(3, 90)
4	$42 + 7(4) + 3(4)^2 = 42 + 28 + 48 = 118$	(4, 118)
5	$42 + 7(5) + 3(5)^2 = 42 + 35 + 75 = 152$	(5, 152)

So, our **5 shares** are:

$$(1, 52), (2, 68), (3, 90), (4, 118), (5, 152)$$

Example of Secret Sharing Techniques

Step 2: Reconstruct the Secret

We need at least 3 shares to recover $S = 42$. Let's use $(1, 52), (2, 68), (3, 90)$.

To recover the secret, we use **Lagrange Interpolation**, which reconstructs $f(0)$ (i.e., the constant term $a_0 = 42$).

The Lagrange basis polynomials are:

$$L_i(0) = \prod_{j \neq i} \frac{x - x_j}{x_i - x_j}$$

Using $x_1 = 1, x_2 = 2, x_3 = 3$:

$$S = f(0) = y_1 L_1(0) + y_2 L_2(0) + y_3 L_3(0)$$

Example of Secret Sharing Techniques

Calculating step-by-step:

$$L_1(0) = \frac{(0 - 2)(0 - 3)}{(1 - 2)(1 - 3)} = \frac{(0 - 2)(0 - 3)}{(1 - 2)(1 - 3)} = \frac{(0 - 2)(0 - 3)}{(1 - 2)(1 - 3)} = \frac{6}{2} = 3$$

$$L_2(0) = \frac{(0 - 1)(0 - 3)}{(2 - 1)(2 - 3)} = \frac{(0 - 1)(0 - 3)}{(2 - 1)(2 - 3)} = \frac{3}{-1} = -3$$

$$L_3(0) = \frac{(0 - 1)(0 - 2)}{(3 - 1)(3 - 2)} = \frac{(0 - 1)(0 - 2)}{(3 - 1)(3 - 2)} = \frac{2}{2} = 1$$

Now, compute S :

$$\begin{aligned} S &= 52(3) + 68(-3) + 90(1) \\ &= 156 - 204 + 90 \\ &= 42 \end{aligned}$$

Thus, we successfully recovered the secret $S = 42$.

Additive Secret Sharing

Additive Secret Sharing

- Each client splits its update into **random additive shares** and sends them to different clients.
- The sum of all shares reconstructs the original update at the server.
- Used in **Secure Multi-Party Computation (MPC)** protocols like **Secure Aggregation** by Bonawitz et al. (Google FL).
- **Advantages:** Computationally efficient, supports dropout-resilient FL.
- **Challenges:** Requires communication between clients, handling dropout clients increases complexity.

Additive Secret Sharing

Each client C_i splits its update x_i into multiple **random shares** so that no single share reveals any information about x_i . Only the **sum of all shares** can reveal the true value.

Steps in Additive Secret Sharing

1. Secret Splitting (at each client)

- Each client C_i has a local update x_i .
- It generates **random shares** such that:

$$x_i = s_{i,1} + s_{i,2} + \cdots + s_{i,N} \quad \text{mod } P$$

where:

where P is a prime number or a large modulus used to prevent overflow.

- $s_{i,1}, s_{i,2}, \dots, s_{i,N-1}$ are random values.
- $s_{i,N}$ is computed so that the sum equals x_i (modulo a prime P).

Additive Secret Sharing

2. Share Distribution

- Each client C_i **sends shares** to different clients:
 - $s_{i,j}$ is sent to client C_j (for all $j \neq i$).
 - Each client **receives shares** from other clients.

3. Aggregation (at the Server)

- Each client sums up the received shares and sends the result to the server:

$$S_j = \sum_{i=1}^N s_{i,j}$$

- The server **adds up all client sums**:

$$\sum_{j=1}^N S_j = \sum_{i=1}^N x_i$$

- Since all intermediate randomness cancels out, the server retrieves the **sum of updates** but not the individual updates.

Additive Secret Sharing

Let's say a model update (e.g., a weight value) is **42**, and we want to split it among three parties.

1. Generate two random numbers:

- $S_1 = 10$
- $S_2 = 25$

2. Compute the third share:

$$\bullet \quad S_3 = 42 - (10 + 25) = 7$$

3. Distribute shares:

- Party 1 receives $S_1 = 10$
- Party 2 receives $S_2 = 25$
- Party 3 receives $S_3 = 7$

4. Each party processes computations on its share.

5. To reconstruct:

$$\bullet \quad 10 + 25 + 7 = 42$$

Verifiable Secret Sharing (VSS)

- Verifiable Secret Sharing (VSS) is an enhanced version of **Shamir's Secret Sharing (SSS)** that includes a verification mechanism to prevent malicious participants from providing incorrect shares.

Verifiable Secret Sharing (VSS)

Why is VSS Needed?

While **Shamir's Secret Sharing (SSS)** ensures that a secret is securely divided and only reconstructable with a threshold number of shares, it **does not** verify whether:

- ✓ The dealer (the entity distributing the shares) is honest.
- ✓ The shares provided by participants during reconstruction are correct.

Problem Example:

If a malicious party injects incorrect shares into the reconstruction phase, the recovered secret will be wrong.

VSS solves this by adding a verification mechanism that allows participants to **validate their shares** without revealing the secret.

Verifiable Secret Sharing (VSS)

1. Secret Splitting (Like SSS)

- The secret S is encoded in a **random polynomial** $f(x)$.
- Each participant P_i receives a share (x_i, y_i) .

2. Public Commitments

- The dealer publicly commits to the coefficients of the polynomial using a cryptographic commitment scheme (e.g., **Pedersen Commitments** or **Feldman's VSS**).

3. Verification of Shares

- Participants can verify whether their received shares are consistent with the public commitments without learning the secret.
- If a participant finds an incorrect share, they can reject it before reconstruction.

4. Secret Reconstruction

- Only valid shares participate in reconstructing S , ensuring correctness.

Verifiable Secret Sharing (VSS)

Types of VSS

1. Feldman's VSS

- Uses **public-key cryptography** to commit to polynomial coefficients.
- Participants verify shares using modular exponentiation.

2. Pedersen's VSS

- Uses **homomorphic commitments** to add an extra layer of secrecy, preventing attackers from inferring secret polynomial coefficients.

Example of Verifiable Secret Sharing (VSS)

Step 3: Compute Public Commitments

We use **modular arithmetic** for verification.

- Choose a **large prime number** $p = 97$.
- Select a **generator** $g = 5$ of a cyclic group.

Compute commitments for each polynomial coefficient:

$$C_0 = g^{a_0} \pmod{p} = 5^{42} \pmod{97} = 54$$

$$C_1 = g^{a_1} \pmod{p} = 5^7 \pmod{97} = 78$$

$$C_2 = g^{a_2} \pmod{p} = 5^3 \pmod{97} = 125 \pmod{97} = 28$$

Thus, our **public commitments** are:

$$C_0 = 54, \quad C_1 = 78, \quad C_2 = 28$$

Example of Verifiable Secret Sharing (VSS)

Step 4: Verify Shares

Each participant verifies their share (x, y) using the commitment equation:

$$g^{y_i} \stackrel{?}{=} C_0 \cdot C_1^{x_i} \cdot C_2^{x_i^2} \pmod{p}$$

Example Verification for $P_1 = (1, 52)$

$$g^{52} \pmod{97} = 5^{52} \pmod{97} = 40$$

Now, compute the right-hand side:

$$\begin{aligned} & C_0 \cdot C_1^1 \cdot C_2^{1^2} \pmod{97} \\ &= (54 \cdot 78^1 \cdot 28^1) \pmod{97} \\ &= (54 \cdot 78 \cdot 28) \pmod{97} \\ &= 117936 \pmod{97} = 40 \end{aligned}$$

Since both sides are equal:

$$40 = 40$$

P_1 's share is valid .

Example of Verifiable Secret Sharing (VSS)

Step 5: Secret Reconstruction (Using Lagrange Interpolation)

To reconstruct the secret, we use **any 3 valid shares** and apply **Lagrange Interpolation** (as done in SSS). This will recover:

$$S = 42$$

Thus, VSS ensures that **only valid shares participate in secret reconstruction**.